

1 Overview

1.1 Location \$(AMDAPPSDKSAMPLESROOT)\samples\C++Amp\examples

1.2 How to Run See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The sample executable is in \$(AMDAPPSDKSAMPLESROOT)\samples\C++Amp\bin\x86\ for 32-bit builds, and \$(AMDAPPSDKSAMPLESROOT)\samples\C++Amp\bin\x86_64\ for 64-bit builds.

Type the following command(s).

1. GEMM

This runs the program with the default options: -i 10 -s 2048

2. GEMM -h

This prints the help file.

Ensure you have installed Microsoft® Visual Studio® 2012 or higher.

1.3 Command Line Options Table 1 lists, and briefly describes, the command line options.

Table 1 Command Line Options

Short Form	Long Form	Description
-h	--help	Show all command options and their respective meaning.
-q	--quiet	Quiet mode. Suppresses text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing.
-d	--deviceId	Select deviceId to be used (0 to N-1, where N is the ID of the device to be used).
-v	--version	AMD APP SDK version string.
-2	--size	Edge length of square matrices A, B, and C.
-i	--iterations	Number of times to repeat each algorithm.
-V	--arrayview	Use array_view instead of array.

2 Introduction

The General Matrix Multiply (GEMM) is a subroutine in the Basic Linear Algebra Subprograms (BLAS). It performs matrix multiplication: the multiplication of two matrices. GEMM is often tuned by high-performance computing (HPC) vendors to run as fast as possible.

The GEMM routine calculates the new value of matrix C based on the matrix-product of matrices A and B, as well as the old value of matrix C.

$$.C = aAB + \beta C$$

where a and β values are scalar coefficients.

According to the data type, GEMM has four versions: SGEMM, DGEMM, CGEMM, and ZGEMM for handling float, double, complex, and complex-double data types, respectively. This sample implements SGEMM NN/TN and DGEMM NN/TN. N means normal; T means transpose.

3 References

1. Dongarra, Jack J.; Croz, Jeremy Du; Hammarling, Sven; Duff, Iain S. (1990), "A set of level 3 basic linear algebra subprograms", [ACM Transactions on Mathematical Software](#) 16 (1): 1–17, doi:10.1145/77626.79170, ISSN 0098-3500.
2. Golub, Gene H.; Van Loan, Charles F. (1996), Matrix Computations (3rd ed.), Johns Hopkins, ISBN 978-0-8018-5414-9.
3. GSL Team (2007), "[§12.1.3 Level 3 GSL BLAS Interface](#)", GNU Scientific Library. Reference Manual, http://www.gnu.org/software/gsl/manual/html_node/Level-3-GSL-BLAS-Interface.html.
4. Goto, Kazushige; van de Geijn, Robert A. (2008), "Anatomy of High-Performance Matrix Multiplication", [ACM Transactions on Mathematical Software](#) 34 (3): Article 12, 25 pages, doi:10.1145/1356052.1356053.

Contact

Advanced Micro Devices, Inc.
One AMD Place
P.O. Box 3453
Sunnyvale, CA, 94088-3453
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:
URL: developer.amd.com/appsdk
Developing: developer.amd.com/
Support: developer.amd.com/appsdksupport
Forum: developer.amd.com/openclforum



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

Copyright and Trademarks

© 2012 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.