

## 1 Overview

**1.1 Location** \$(AMDAPPSDKSAMPLESROOT)\samples\C++Amp\examples

**1.2 How to Run** See the *Getting Started* guide for how to build samples. You first must compile the sample.

Use the command line to change to the directory where the executable is located. The default executables are placed in \$(AMDAPPSDKSAMPLESROOT)\samples\C++Amp\bin\x86\ for 32-bit builds, and \$(AMDAPPSDKSAMPLESROOT)\samples\C++Amp\bin\x86\_64\ for 64-bit builds.

Type the following command(s).

1. SPMV

This runs the program with the default options -l 16384

2. SPMV -h

This prints the help file.

Ensure Microsoft® Visual Studio® 2012 or higher is installed.

**1.3 Command Line Options** Table 1 lists, and briefly describes, the command line options.

**Table 1 Command Line Options**

Short Form	Long Form	Description
-h	--help	Show all command options and their respective meaning.
-q	--quiet	Quiet mode. Suppresses text output.
-e	--verify	Verify results against reference implementation.
-t	--timing	Print timing-related statistics.
-v	--version	AMD APP SDK version string.
-d	--deviceId	Select deviceId to be used (0 to N-1, where N is the number of available devices).
-i	--iterations	Number of times to repeat each algorithm.
-V	--array_view	Use array_view instead of array.
-l	--length	Length of row/column of square matrix.
-p	--ellpackr	Use Ellpackr format instead of Compressed Sparse Row (CSR) format.


## 2 Introduction

This sample computes large sparse matrix-vector multiplication on the GPU in two different storage formats: CSR or ELLPACKR.

## 3 Implementation Details

Sparse matrix usually refers to a class of matrices, the majority of their position are zeros. For Sparse matrix, in order to save storage space, we usually save only those non-zero elements, and record their position, we do not record the position of the elements that are zeros.

The CSR (Compressed Sparse Row) format is shown in [Figure 1](#). The VAL array stores all the non-zero elements of the matrix. The J array has the same number of elements as the VAL array, and it stores the column index of the corresponding element in VAL. The I array stores the index in array J where a new row starts.



**Figure 1 CSR Format**

We can assign the computation of each row to a thread, so the total number of thread is equal to the number of rows.

The SPMV algorithm of CSR format is composed of:

1. Compute the matrix row our thread is responsible for.
2. Compute start/end indices of our matrix row in the sparse value array.
3. Build a partial sum for the row, striding by cCSRNTthreadsPerRow.

```
for(int i = colStart + localIdxInRow; i < colEnd; i += cCSRNTthreadsPerRow)
{
    thrPartialSum += csrVals_view[i] * inVector_view[csrCols_view[i]];
}
```

4. Reduce partial sums across the tile. The code implementation is as follow:

```
if(localIdxInRow < 16) partialSums[idx.local[0]] += partialSums[idx.local[0] + 16];
idx.barrier.wait_with_tile_static_memory_fence();
if(localIdxInRow < 8) partialSums[idx.local[0]] += partialSums[idx.local[0] + 8];
idx.barrier.wait_with_tile_static_memory_fence();
if(localIdxInRow < 4) partialSums[idx.local[0]] += partialSums[idx.local[0] + 4];
idx.barrier.wait_with_tile_static_memory_fence();
if(localIdxInRow < 2) partialSums[idx.local[0]] += partialSums[idx.local[0] + 2];
idx.barrier.wait_with_tile_static_memory_fence();
if(localIdxInRow == 0)
{
    outVector[rowIdx] = partialSums[idx.local[0]] + partialSums[idx.local[0] + 1];
}
```

ELLPACKR (Ellpack-ltpack) format uses a value matrix, *V*, that stores all non-zero values in a matrix, and each row is zero-padded to the same size as the row with the maximum number of non-zero elements. The row index in the *V* matrix is the same as in the original sparse matrix. The column index of each value is stored in the matrix *J*. The significant computation is:

```
for(int i = 0; i < rowLength; ++i)
{
    const int offset = i * denseEdge + idx[0];
    sum += ellpackrVals[offset] * inVector[ellpackrCols[offset]];
}
```

---

**Contact**

Advanced Micro Devices, Inc.  
One AMD Place  
P.O. Box 3453  
Sunnyvale, CA, 94088-3453  
Phone: +1.408.749.4000

For AMD Accelerated Parallel Processing:  
URL: [developer.amd.com/appsdk](http://developer.amd.com/appsdk)  
Developing: [developer.amd.com/](http://developer.amd.com/)  
Support: [developer.amd.com/appsdksupport](http://developer.amd.com/appsdksupport)  
Forum: [developer.amd.com/openclforum](http://developer.amd.com/openclforum)



The contents of this document are provided in connection with Advanced Micro Devices, Inc. ("AMD") products. AMD makes no representations or warranties with respect to the accuracy or completeness of the contents of this publication and reserves the right to make changes to specifications and product descriptions at any time without notice. The information contained herein may be of a preliminary or advance nature and is subject to change without notice. No license, whether express, implied, arising by estoppel or otherwise, to any intellectual property rights is granted by this publication. Except as set forth in AMD's Standard Terms and Conditions of Sale, AMD assumes no liability whatsoever, and disclaims any express or implied warranty, relating to its products including, but not limited to, the implied warranty of merchantability, fitness for a particular purpose, or infringement of any intellectual property right.

AMD's products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD's product could create a situation where personal injury, death, or severe property or environmental damage may occur. AMD reserves the right to discontinue or make changes to its products at any time without notice.

**Copyright and Trademarks**

© 2012 Advanced Micro Devices, Inc. All rights reserved. AMD, the AMD Arrow logo, ATI, the ATI logo, Radeon, FireStream, and combinations thereof are trademarks of Advanced Micro Devices, Inc. OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos. Other names are for informational purposes only and may be trademarks of their respective owners.